

PL4DS Meeting

Feb 23rd 2018

Guido Salvaneschi

TU Darmstadt - Germany

PhD - Politecnico di Milano, Italy

PL for (distributed) adaptive systems

Erlang: Hot swap of code modules



(Distributed) Reactive Programming

REScala programming language

Distribution

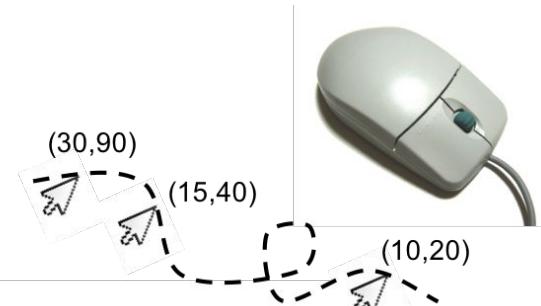
Concurrency

Debugging

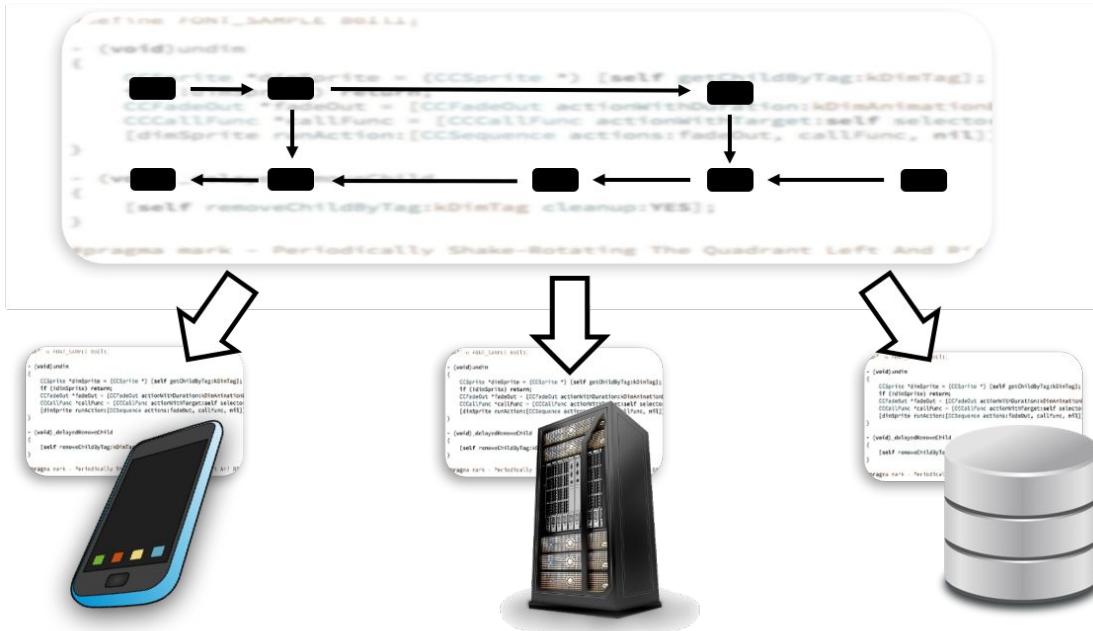
Controlled experiments

```
val position: Signal[(Int,Int)] = mouse.position  
val shifted: Signal[(Int,Int)] = Signal{ mouse.position + (10,10) }
```

```
evt clicked: Event[Unit] = mouse.clicked  
val lastClick: Signal[(Int,Int)] = position snapshot clicked
```



Tierless Programming: ScalaLoci



Data Streams cross tier boundaries

Compiler splits the code

ScalaLoci

Placement types: Server[Int], Client[Bool]

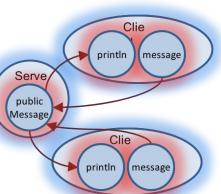
```
import architecture.MultiClientServer._  
import serializable.upickle._  
import rescalaTransmitter._
```

```
@multitier
object Chat {
    class Server extends ServerPeer[Client] {
        def connect = TCP(port = 43053) }
    class Client extends ClientPeer[Server] {
        def connect = load (TCP) ("client.config") }
```

```
val message = placed[Client] { Evt[String] }
```

```
val publicMessage = placed[Server] {  
    message.asLocalSeq map { case ( , message) => message } }  
}
```

```
placed[Client].main {
    publicMessage.asLocal observe println
    for (line <- io.Source.stdin.getLines)
        message fire line } }
```



```
val ballSize = 20
val maxX = 500
val maxY = 400
val leftPos = 30
val rightPos = 770
val initPosition = Point(400, 200)
val initSpeed = Point(10, 0)
```

```
    val ball: Signal[Point] = tick.fold(initPosition) {  
      (ball, tick) => ball + Vector(0, 1)  
    }  
  }
```

```
trait Server extends ServerPeer[Client]
trait Client extends ClientPeer[Server]

val ballSize = 20
val maxx = 800
val maxy = 400
val leftPos = 30
val rightPos = 770
val initPosition = Point(400, 200)
val initSpeed = Point(10, 8)
```

```

val clientHouseY = placed[Client].y
Signal { v1, housePosition(v1.y) }

val isPlaying = placed[Server].loc.y {
    Signal { remote[Client].connected, size > 2 }
}

val ball: Signal = InServer = placed {
    tick.foldInInitPosition() { (ball, t) =>
        if (!isPlaying) ball + speed.get * t
    }
}

```

Carla Ferreira

Universidade NOVA de Lisboa

PhD - University of Southampton, UK



Formal methods

Concurrency and distribution

Calculi for (failure recovery of) long running transactions

Principled methods and verification techniques for cloud applications

Challenges in globally distributed data

- To ensure availability, cloud applications tend to trade strong consistency for weaker models
- These weaker models do not ensure global ADT invariants
 - Solution: Strengthen consistency selectively
- Large cloud providers started to provide multiple consistency choices
- Hard to figure out the **minimum consistency** necessary to **maintain ADT invariants**

CISE consistency model

- A generic model that expresses most existing consistency models
 - Each operation acquires a set of tokens
 - Operations with conflicting tokens cannot run concurrently
- **First proof rule** to check correctness of weakly consistent applications
 - Based on rely-guarantee reasoning
 - Polynomial time analysis
 - Assumes causal consistency
- Static analysis tool (Boogie)

CISE analysis

- Allows developers to build applications assuming initially a sequential setting
 - Then, the analysis detects what are the problematic operations if executed concurrently
- Specification effort is high: input is a Boogie spec of the application
- Cloud apps moved from monoliths to microservices
 - Extend the proof rule to support compositional reasoning

Carlos Baquero

Universidade do Minho - Portugal

HASLab - INESC TEC

(Large Scale) Distributed Systems

Eventual Consistency - CRDTs

Distributed Data Aggregation



Eventual Consistency ADT Design

Available under Partitions designs (CAP)

Preservation of Sequential Semantics

Design choices under concurrency

Add vs Remove Wins

Outcomes not explained by sequential executions

Multi-value Registers

Causal Consistency

Eventual Consistency Support

Reliable FIFO Channels

Preserve source FIFO

Trees provide Causal Consistency

Reliable Causal Delivery

Causal Consistency for arbitrary graphs

Delivery pre-conditions vs batch propagation

FIFO not across TCP connection incarnations

Crossing the CAP Barrier

“Observable” Causal Consistency strongest under AP

Towards global invariants:

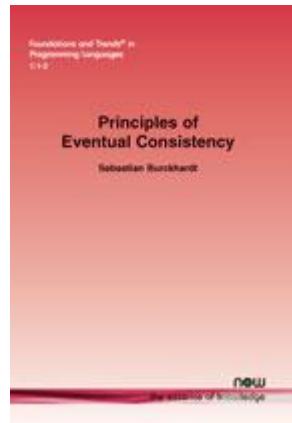
- A) Escrow, Tokens, Reliable transfers
- B) Consensus Services

Sebastian Burckhardt

Microsoft Research, Redmond

I like to make complicated things look simple.

Theoretical Interest: Specification/Verification/Optimality of
~~ConcurrentDistributed MemoryObjectsServices~~



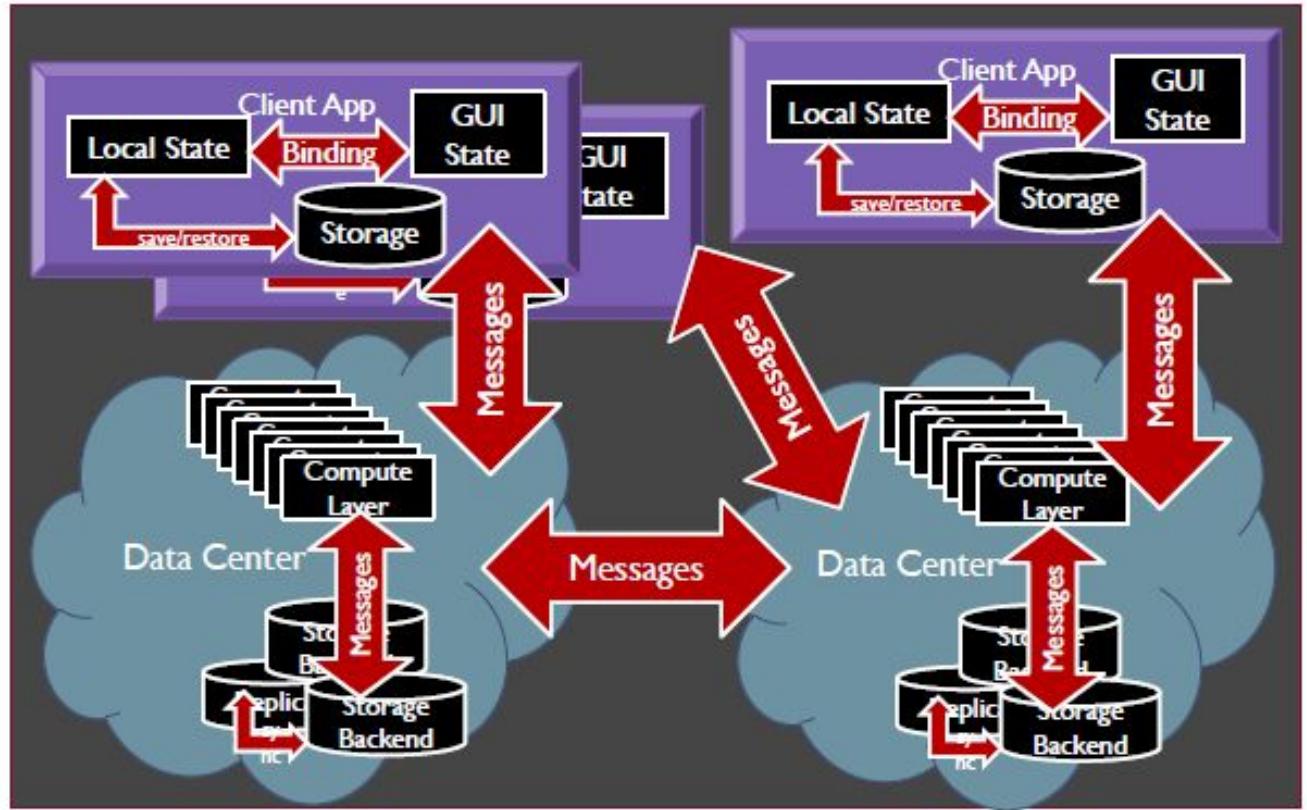
Practical Interest: figuring out the new abstraction stack

(how to build applications in the age of devices+services)

Reactivity, streaming, cloud storage, browsers...

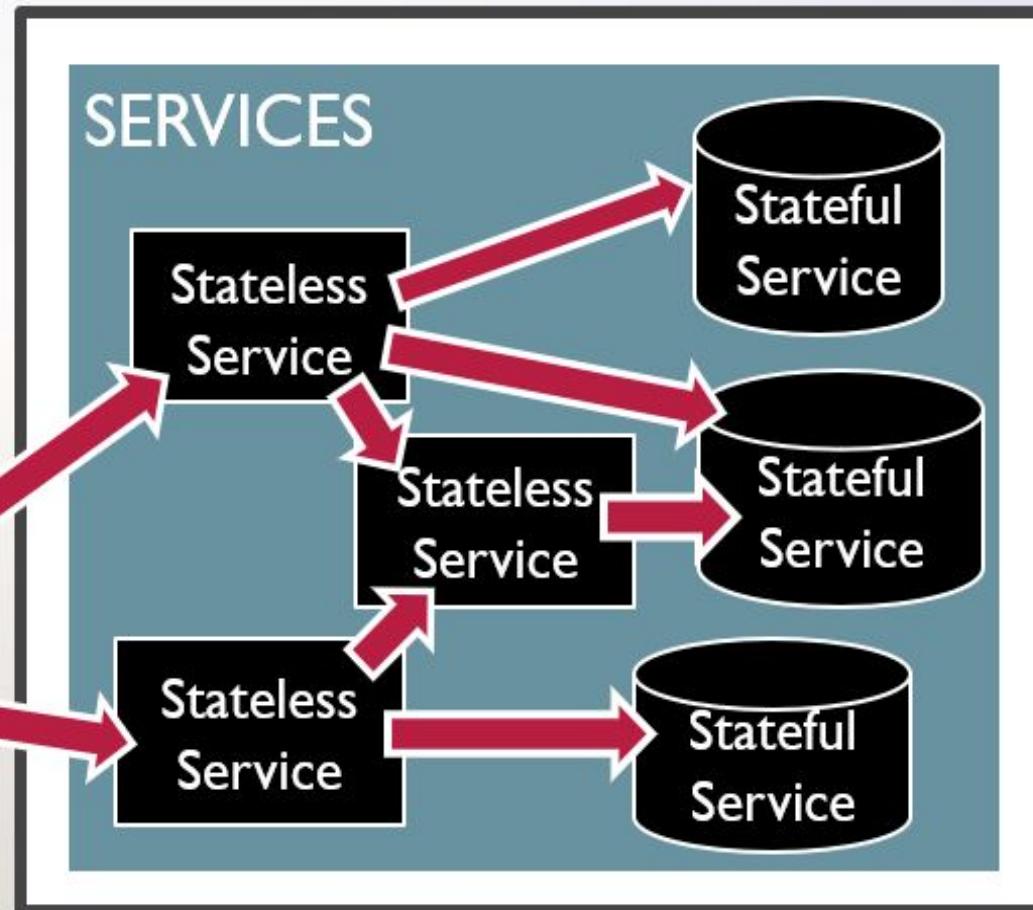
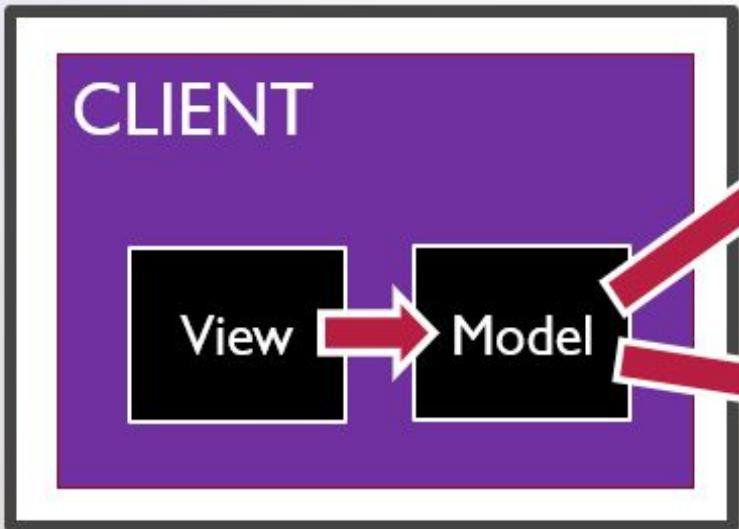
PERFECT STORM OF COMPLEXITY

- security,
concurrency,
parallelism,
distribution,
consistency,
failures,
dynamic topology.



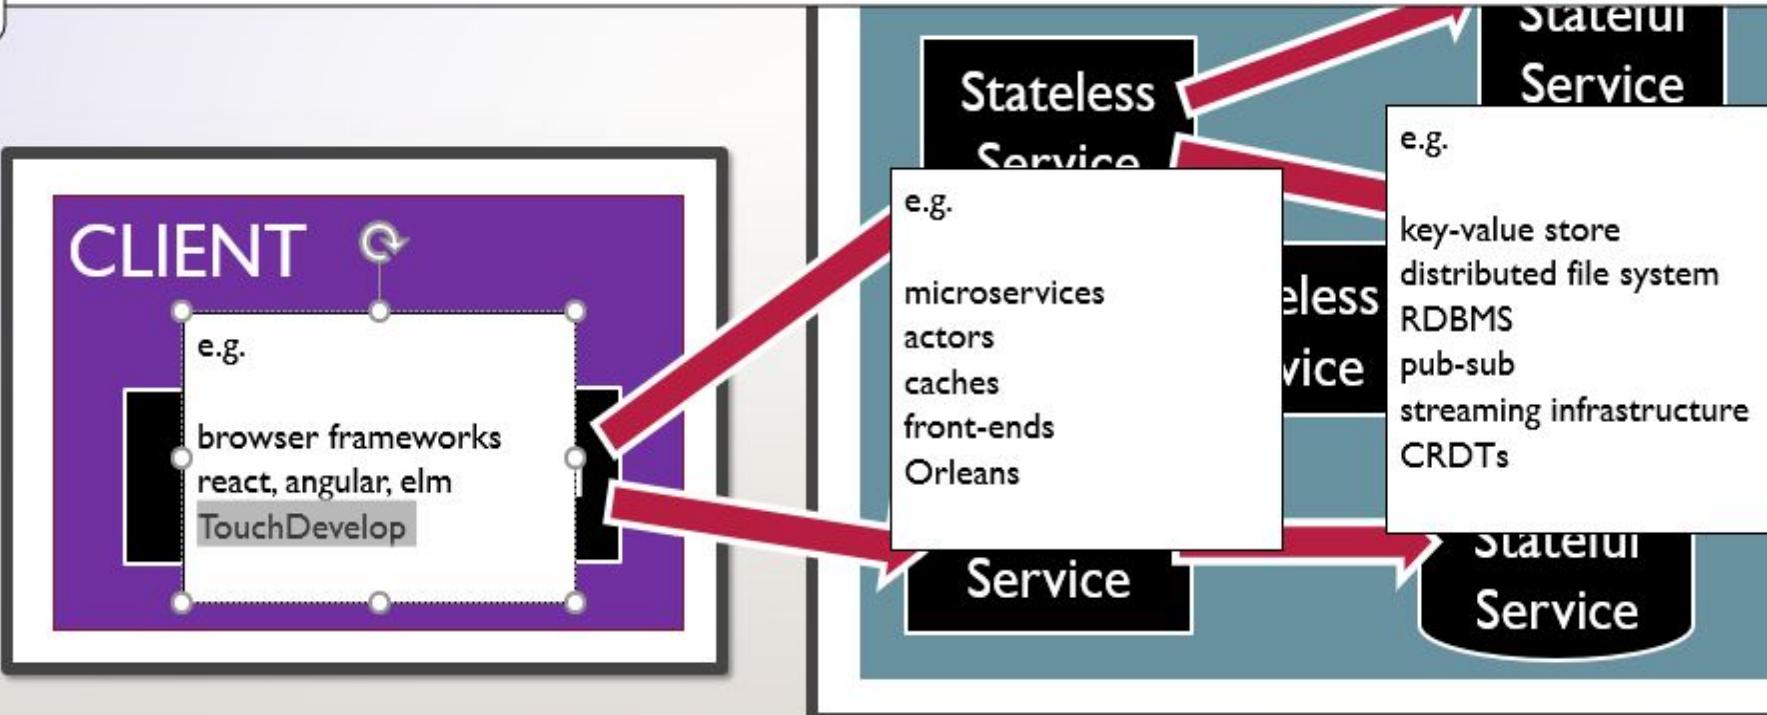
EMERGING PATTERN

COMPOSED SERVICES



application specific
runs application code
unreliable, unavailable

reliable, available
heavily engineered
general purpose
no application code*



Damien Zufferey

- PhD: IST Austria (2013)
- Postdoc: MIT CSAIL (2013-2016)
- Research group leader: MPI-SWS (2016-...)

Research related to distributed systems:

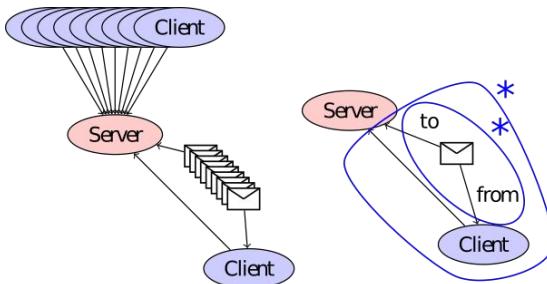
- Verification of mobile systems (pi-calculus)
- Programming abstractions for messages and faults
- Coordination for CPS

More info at <http://dzufferey.github.io/>



Damien Zufferey: Research in Pictures

Verification of mobile processes

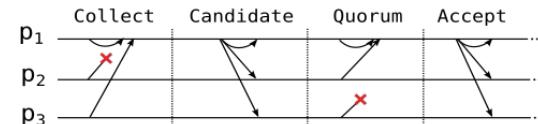


- Process creation
- Changing communication topology
- ...

Fault and messages:

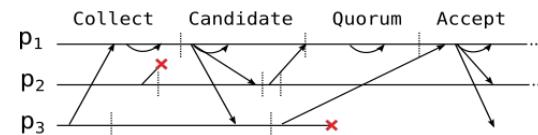
Using communication-closure to simplify programming and verification

Lockstep Semantics



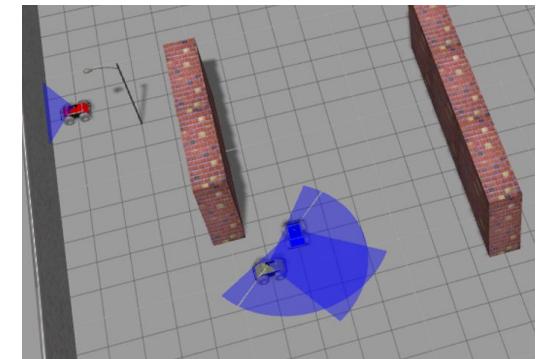
↑ Indistinguishable

Asynchronous Execution



Coordination for CPS

When messages carries informations about the “real” world



Aleksandar Prokopec

- 2009-2014 PhD at EPFL
 - Scala Team
- 2014-2016 Google
 - Google Maps
- 2016-2018 Principal Researcher at Oracle Labs
 - GraalVM

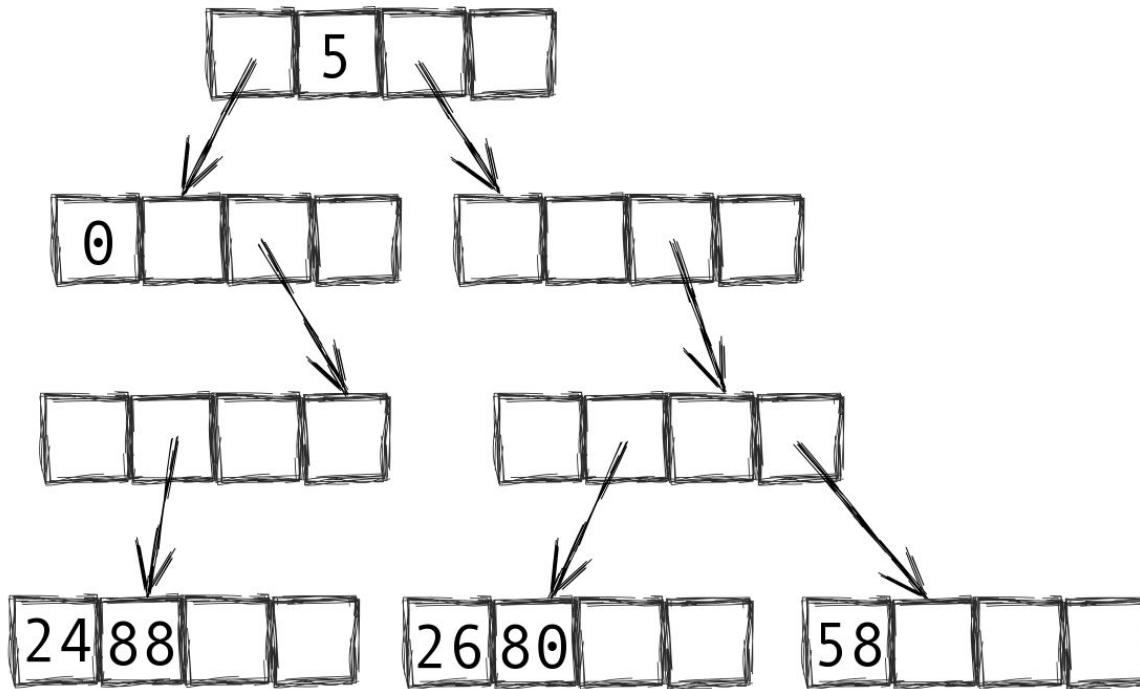
Research:

- concurrency, non-blocking data structures
- programming models for distributed systems
- compilers, optimizations, type systems

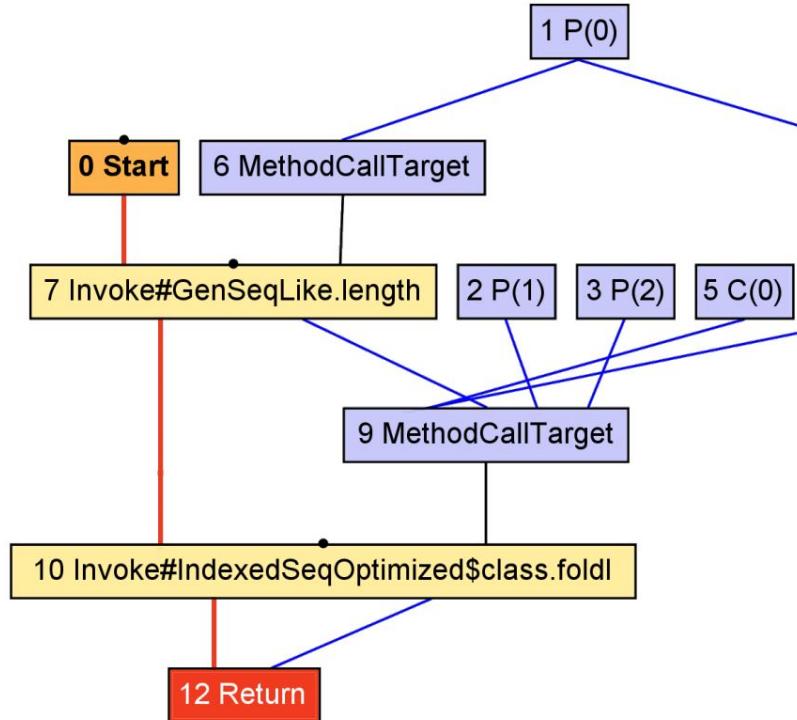


Papers, CV, other details: <http://aleksandar-prokopec.com>

Aleksandar Prokopec: Concurrent Data Structures

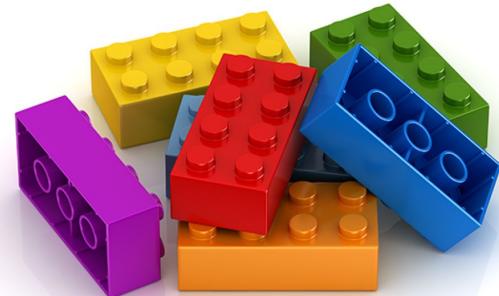


Aleksandar Prokopec: Compilers



Aleksandar Prokopec: Distributed Computing

```
process {
    val work = broadcast(workers)
    val decisions = paxos(leaders)
    while (true) {
        work ! receive(decisions)
    }
}
```



Wolfgang De Meuter

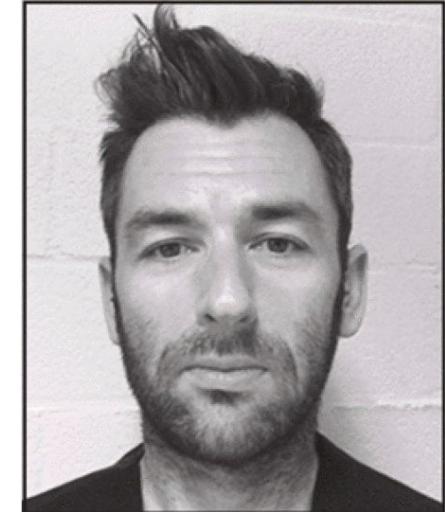
Prof @ Vrije Universiteit Brussel (Software Languages Lab)

Topics of my team:

- Distributed & concurrent programming languages
Reactive Programming languages (procedural, logic)
- Reactive Big Data processing languages

Current Application **Themes** (project portfolio):

- Participatory Sensing (citizen science, crowd sensing)
- Cyber Physical Systems
- Internet of Things
- Security by CEP
- Smart Grids



Technologies:

- Scheme, Clojure, Lisp
- Smalltalk, AmbientTalk
- Spark, Flink,
- JavaScript, TypeScript
- Scala, Haskell, ...

Topic #1: Concurrent & Distributed Languages

- AmbientTalk, AmbientTalk/M, AmbientTalk/R
 - “Failure is the rule rather than the exception”
- Clojure extension with STM + Futures + actors
 - “One paradigm doesn’t fit all”
- Static Analysis of Actor-based languages
 - “Prevent your actor from exploding”
- Reactive Programming for Actor Coordination
 - “Coordination logic = reacting to complex message patterns”



Janwillem
 Swalens



2016



Quentin
Stievenart



2017



Humberto
Avila Rodriguez

Topic #2: Reactive Programming Languages

Combining reactive programming with distribution, multiple external signals, imperativeness/loops (e.g. sequencing, coordination):

The Actor-Reactor model

Fundamentals of RP:

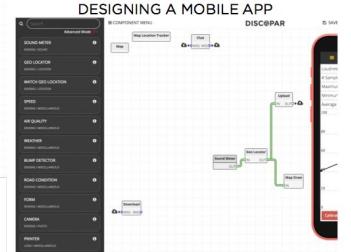
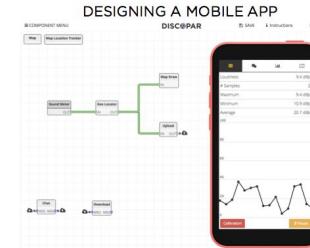
- Reactive VM & instruction set
- Higher Order Reactors 4 Live Programming
- Reactive Meta-programming
- Logic Reactive Programming with RETE



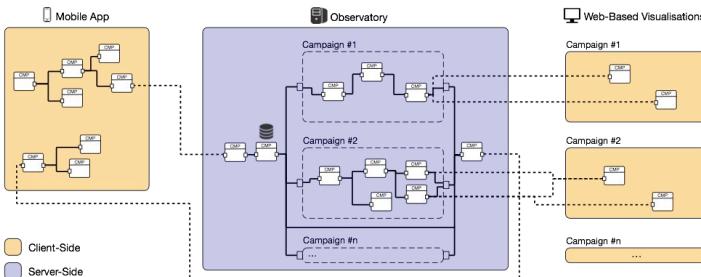
Topic #3: Reactive Big Data Processing Languages



DISCOPAR



SERVER-SIDE DATA PROCESSING



TZU-CHUN (Gina) CHEN

TU Darmstadt - Germany (2014- 2018)

Postdoc, working with Prof. Patrick Eugster
ERC LiveSoft project, DSP group

University of Torino - Italy (2013- 2014)

Postdoc, working with Prof. Mariangiola Dezani
and Prof. Luca Padovani

PhD -- Queen Mary, University of London, UK (2013)

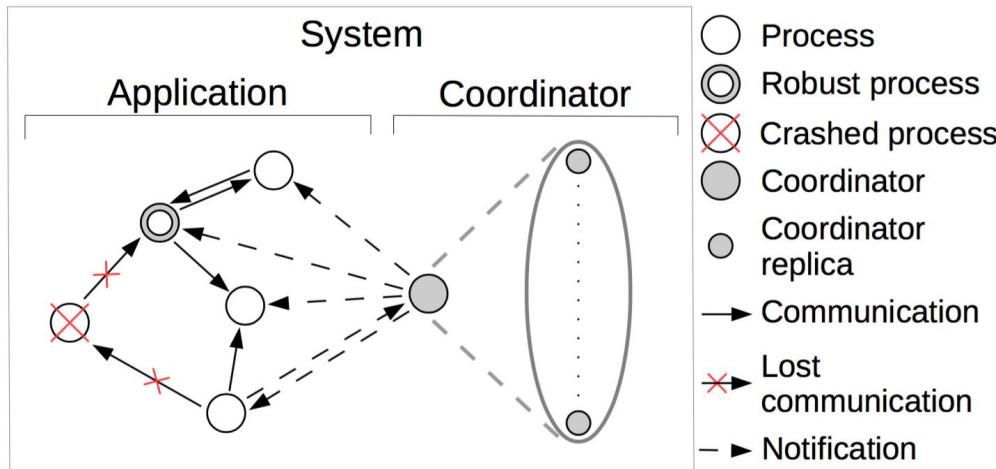
Research interests

Type systems for distributed programming
PL for safe&efficient failure handling in DS/IoT
Safe and secure communication for DS/IoT



Challenges in DS Failure Handling @ N Processes

- Asynchrony, concurrency, failure (process crash)
 - It is *impossible* for processes to reach an agreement
-- by Fischer, Lynch, Paterson JACM' 85
- We introduce a coordinator Ψ to the system for failure handling

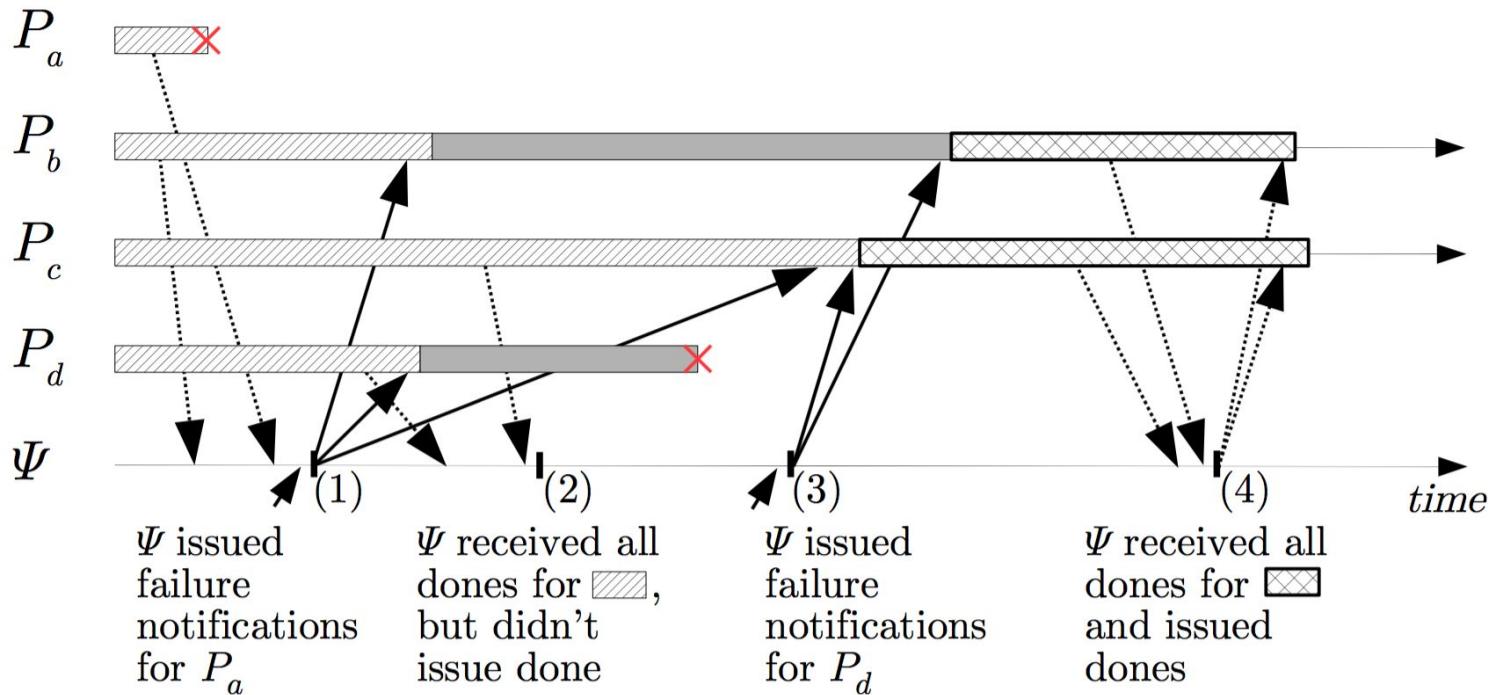


Ensuring communication safety and system progress:

- From PL perspective
- Session-based process calculus with failure handling capability
- Session-type-based type disciplines

Non-trivial DS Failure Handling

Task: (κ, \emptyset) $(\kappa, \{P_a\})$ $(\kappa, \{P_a, P_d\})$



Non-trivial DS Failure Handling

Related Publications

A Type Theory for Robust Failure Handling in Distributed Systems. FORTE 2016 --- no crash, no coordinator model

A Typing Disciplines for Statically Verified Crash Failure Handling in Distributed Systems. ESOP 2018 (accepted) -- with coordinator model

Ongoing works

Formalising Zookeeper for Spark Failure Handling--PhD Malte Viering

Implementing Our Coordinator Model for Spark--PhDMalte Viering

Failure Handling/Safe Communication for IoT Applications--PhD Seema Kumar

Peter Van Roy



Université catholique de Louvain – Belgium

Programming Languages and Distributed Computing

LightKone H2020 project: Lightweight Computation
for Networks at the Edge (2017-2019)

Concepts, Techniques, and Models of Computer
Programming (MIT Press, 2004)

Taking advantage of weak synchronization: convergent computation and
synchronization-free services



Convergent computation

- Computation is always converging to a result
 - Generalizes eventual consistency to simplify programming
 - Automatically maintains correctness of program data with respect to the external world
 - Naturally tolerates node and communication problems, e.g., on edge networks
- Lasp (“Lattice Programming”) lasp-lang.org
 - Programming language for writing convergent computations
 - Basis for an edge computing platform being implemented in the LightKone project

Lasp example



- Sets connected with a map (Erlang syntax):

```
s1=declare(set),  
bind(S1, {add, [1,2,3]}),  
S2=declare(set),  
map(S1, fun(X)->X*X end, S2).
```

- Deterministic dataflow functional semantics
- This is a convergent computation
 - Tolerates node and communication failures
 - Keeps consistency using only weak synchronization
- It has an efficient implementation
 - Graph of CRDTs connected with operations
 - Communication layer based on hybrid gossip



unbalancedparentheses [Follow](#)

Federico Carrone. A happy member of The Erlang, Rust/ML and Lisp Evangelism Strike...
May 9 · 8 min read

Lasp: a little further down the Erlang rabbithole

A few years ago I found Lasp: “*a suite of libraries aimed at providing a comprehensive programming system for planetary scale Elixir and Erlang applications*”. At this point it should come as no surprise for you to learn that here at Not a Monad Tutorial we are interested in distributed systems and Erlang. After playing a little bit with Lasp I watched a few talks by its creator: Christopher Meiklejohn. After watching his talk “*Distributed, Eventually Consistent Computations*” I decided it was time to interview Christopher.

Reach me via twitter at @unbalancedparen if you have any comments or interview request for This is not a Monad tutorial. **Stay tuned!**

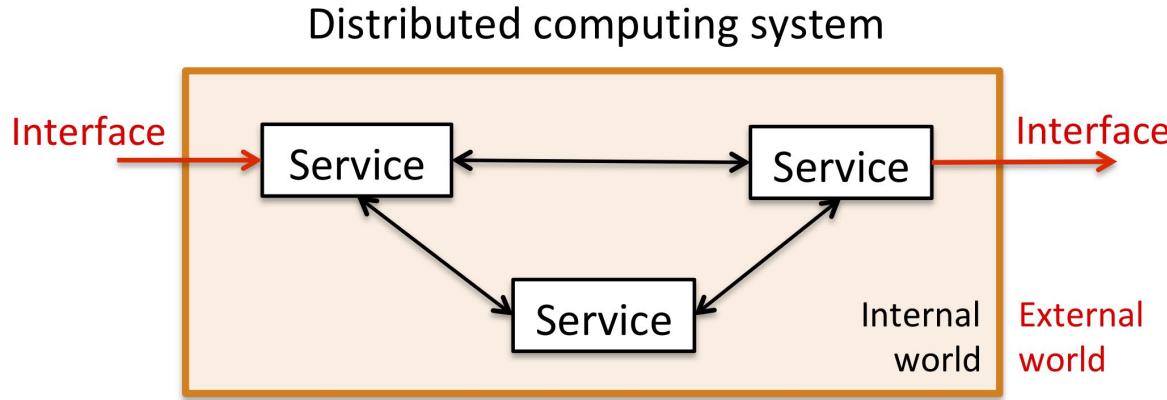
Discuss and vote at lobsters, reddit and hn.



What is Lasp?

Originally, Lasp was a programming model designed for deterministic distributed computing with weak synchronization. Lasp’s

Synchronization-free services



- Most existing services synchronize unnecessarily at their APIs
- Instead, give the *overall system* a synchronization boundary
 - Inside the boundary, *all services use weak synchronization* to interact
 - Strong synchronization is only needed at the boundary, to interface with the external world

Programming Language Frontends

Abstractions

Notations

Analyses

Editors

Sebastian Erdweg



Language Design — Domain-Specific Languages

Corrl: programmable event correlation in CEP

```
with slidingWindow(TVReviews, 1 Month)
correlate {
    release from TVReleases
    reviews(5) from TVReviews
    where
        distinct(reviews)
        forall(reviews)(_.rating >= 3.5)
        forall(reviews)(equal(_.model, release.model))
    yield release
}
```

i3QL: privacy-aware, distributed, incremental LINQ

```
1 val students: Table[Student] = new Table[Student]()
2
3 val sallies: View[Student] =
4 (SELECT (*) FROM students
5 WHERE (s => s.firstName == "Sally")).asMaterialized
6
7 students.add(new Student("Sally", "Fields"))
8 students.add(new Student("George", "Tailor"))
9
10 sallies.foreach(s => println(s.lastName))
11 // prints: "Fields"
12
13 students.add(new Student("Sally", "Joel"))
14 // incremental update of sallies
```

Incremental Program Analysis

6000x speedup

Data flow: Java

0: entry point measure

1: read env

2: ifjump 13

3: read temp

4: map p

5: function call (N) readSensor

6: entry point readSensor

11: unmap p

12: jump 19

nested DFG

call error

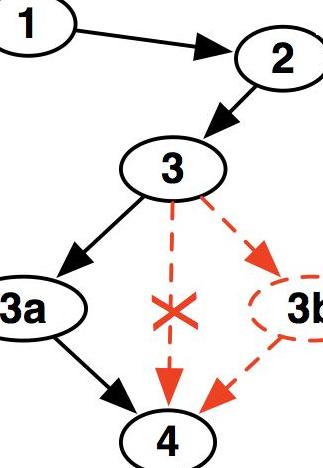
call calibrateEnv

nested DFG

28: end

250x speedup [ASE'16]

Control flow: C, Java



10x speedup

Type checking: Java, Rust

$$\lambda f : \alpha \rightarrow \beta : (\alpha \rightarrow \beta) \rightarrow \alpha \rightarrow \beta | \emptyset$$

$$\lambda x : \alpha : \alpha \rightarrow U_2 | f : \alpha \rightarrow U_2$$

$$app : U_2 | f : U_1 \rightarrow U_2; x : U_1$$

$$f : U_0 | f : U_0$$

$$x : U_1 | x : U_1$$

65x speedup

[OOPSLA'14, ASE'16]

FindBugs: Java



Nobuko Yoshida

<http://mrg.doc.ic.ac.uk>



The screenshot shows the homepage of the Mobility Research Group. At the top, there's a logo with a stylized Greek letter π and the text "session type". Below it, the title "Mobility Research Group" is displayed in a large, serif font. A sub-header "π-calculus, Session Types research at Imperial College" follows. A navigation bar contains links for Home, People, Publications, Grants, Talks, Tutorials, Tools, Awards, and Kohei Honda. The main content area is divided into two sections: "NEWS" on the left and "SELECTED PUBLICATIONS" on the right. The "NEWS" section features a news item about a paper being awarded the ACM SIGPLAN Most Influential POPL Paper Award. The "SELECTED PUBLICATIONS" section lists several papers from 2018, including works by Lange, Ng, Toninho, Yoshida, and others, with links to their abstracts.

Mobility Research Group

π-calculus, Session Types research at Imperial College

Home People Publications Grants Talks Tutorials Tools Awards Kohei Honda

NEWS

The paper *Multiparty asynchronous session types* by Kohei Honda, Nobuko Yoshida, and Marco Carbone, published in POPL 2008 has been awarded the ACM SIGPLAN Most Influential POPL Paper Award today at POPL 2018.

» more

10 Jan 2018

Estafet has published a page on their usage of the Scribble language developed in our group with RedHat and other industry partners.

» more

25 Sep 2017

Nick spoke at Golang UK 2017 on applying behavioural types to verify concurrent Go programs.

SELECTED PUBLICATIONS

2018

Julien Lange , Nicholas Ng , Bernardo Toninho , Nobuko Yoshida : [A Static Verification Framework for Message Passing in Go using Behavioural Types](#). *To appear in ICSE 2018*.

Bernardo Toninho , Nobuko Yoshida : [Depending On Session Typed Process](#). *To appear in FoSSaCS 2018*.

Bernardo Toninho , Nobuko Yoshida : [On Polymorphic Sessions And Functions: A Talk of Two \(Fully Abstract\) Encodings](#). *To appear in ESOP 2018*.

Rumyana Neykova , Raymond Hu , Nobuko Yoshida , Fahd Abdeljallal : [Session Type Providers: Compile-time API Generation for Distributed Protocols with Interaction Refinements in F#](#). *To appear in CC 2018*.



Post-docs:

Simon CASTELLAN

David CASTRO

Francisco FERREIRA

Raymond HU

Rumyana NEYKOVA

Nicholas NG

Alceste SCALAS

PhD Students:

Assel ALTAYEVA

Juliana FRANCO

Eva GRAVERSEN

Interactions with Industries

F#unctional Londoners Meetup Group CC'18

6 days ago · 6:30 PM

Session Types with Fahd Abdeljallal



43 Members

Synopsis: Session types are a formalism to codify the structure of a communication, using types to specify the communication protocol used. This formalism provides the... [LEARN MORE](#)

ECOOP'17

Distributed Systems
vs.
Compositionality

Dr. Roland Kuhn
@rolandkuhn — CTO of Actyx



Current State

- behaviors can be composed both sequentially and concurrently
- effects are not yet tracked
- Scribble generator for Scala not yet there
- theoretical work at Imperial College, London (Prof. Nobuko Yoshida & Alceste Scalas)

ECOOP'16

the morning paper

ICSE'18

an interesting/influential/important paper from the world of CS every weekday morning, as selected by Adrian Colyer

[Home](#) [About](#) [InfoQ QR Editions](#) [Subscribe](#)

A static verification framework for message passing in Go using behavioural types

JANUARY 25, 2018

tags: Concurrency, Programming Languages

[A static verification framework for message passing in Go using behavioural types](#) Lange et al., ICSE 18

With thanks to Alexis Richardson who first forwarded this paper to me.

We're jumping ahead to ICSE 18 now, and a paper that has been accepted for publication there later this year. It fits with the theme we've been exploring this week though, so I thought I'd cover it now. We've seen verification techniques applied in the context of [Rust](#) and [JavaScript](#), looked at the integration of [linear types in Haskell](#), and today it is the turn of Go!

SUBSCRIBE



never miss an issue! The Morning Paper delivered straight to your inbox.

SEARCH

type and press enter

ARCHIVES

Select Month

MOST READ IN THE LAST FEW DAYS

Selected Publications 2017/2018

- ▶ [CC'18] Rumyana Neykova , Raymond Hu, NY, Fahd Abdeljallal: Session Type Providers: Compile-time API Generation for Distributed Protocols with Interaction Refinements in F#.
- ▶ [FoSSaCS'18] Bernardo Toninho, NY: Depending On Session Typed Process.
- ▶ [ESOP'18] Bernardo Toninho, NY: On Polymorphic Sessions And Functions: A Talk of Two (Fully Abstract) Encodings.
- ▶ [ESOP'18] Malte Viering, Tzu-Chun Chen, Patrick Eugster, Raymond Hu , Lukasz Ziarek: A Typing Discipline for Statically Verified Crash Failure Handling in Distributed Systems.
- ▶ [ICSE'18] Julien Lange, Nicholas Ng, Bernardo Toninho, NY : A Static Verification Framework for Message Passing in Go using Behavioural Types.
- ▶ [ECOOP'17] Alceste Scala, Raymond Hu, Ornella Darda, NY: A Linear Decomposition of Multiparty Sessions for Safe Distributed Programming.
- ▶ [COORDINATION'17] Keigo Imai, NY, Shoji Yuen: Session-ocaml: a session-based library with polarities and lenses.
- ▶ [FoSSaCS'17] Julien Lange, NY: On the Undecidability of Asynchronous Session Subtyping.
- ▶ [FASE'17] Raymond Hu, NY: Explicit Connection Actions in Multiparty Session Types.
- ▶ [CC'17] Rumyana Neykova, NY: Let It Recover: Multiparty Protocol-Induced Recovery.
- ▶ [POPL'17] Julien Lange, Nicholas Ng, Bernardo Toninho, NY: Fencing off Go: Liveness and Safety for Channel-based Programming.

Gustavo Petri

IRIF - Paris Diderot - Paris 7

- Verification for weak consistency
- Language-based techniques for distributed applications



PhD. INRIA - Sophia Antipolis

- True concurrency for weak memory models

PostDoc(s)

DePaul University (Chicago): Semantic models for weak memory

Purdue University (Indiana): Verification of concurrent algorithms

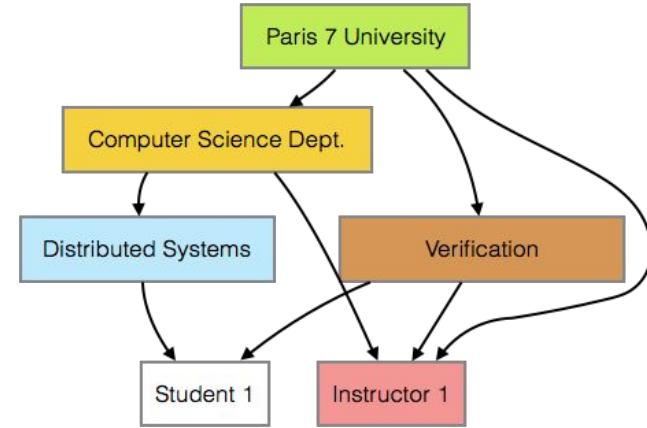
Programming Models for the Cloud

Strong semantics for actor-based transactions

- Dynamic ownership model for transactions
- Provably Linearizable
- Do we really need linearizability?

Programmable Elasticity

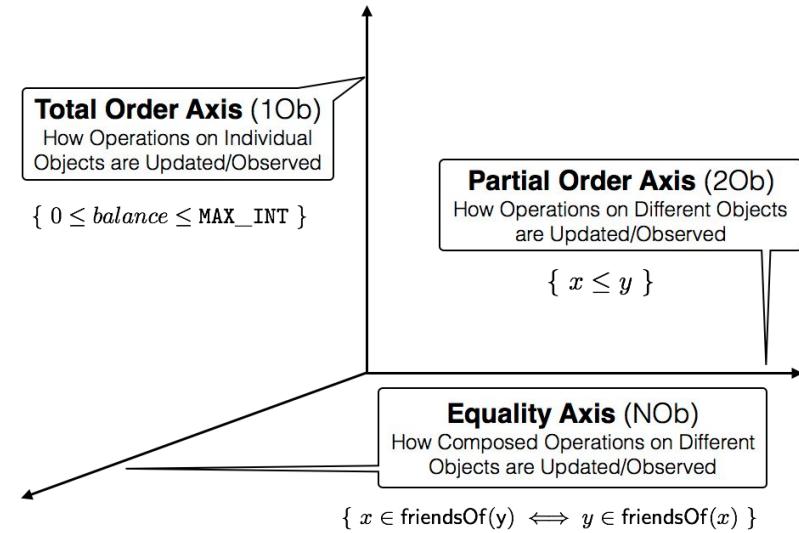
- Let the programmer help decide when and what to migrate across the DC



Verification of Applications under Weak Consistency

Dimensions of Consistency (often times invariants indicate the necessary consistency)

- Strong Consistency =>
General Invariants
- Causal Consistency =>
Referential Integrity, Inequalities
- Transactions =>
Equivalence Relations



Invariants for applications using CRDTs

- What can we prove of applications using many CRDTs and/or transactions
- Verification-aided development

PL4DS – Raymond Hu

- ▶ RA at Imperial College London (in the group of Nobuko Yoshida)
 - ▶ PhD – Imperial College London 2011
- ▶ Implementations and applications of multiparty *session types* (MPST)
 - ▶ Formal languages for message passing between concurrent processes
 - ▶ Static type system for “communication safety”
 - ▶ E.g., freedom from reception errors, deadlocks and orphan messages
- ▶ Practical session types for “mainstream” engineering languages
 - ▶ Language extensions, DSLs, runtime monitoring, code generation
 - ▶ Static typing, dynamic verification, “hybrid” approaches

Scribble: protocol-specific API generation from session types

- ▶ MPST-based toolchain for Java (also .NET, Go, ...)
 - ▶ Protocol validation: MPST syntax + model checking (+ SAT solving)
 - ▶ Endpoint programming using generated distributed APIs

The image shows two side-by-side screenshots. On the left is a screenshot of the RFC 5321 page on the IETF website, titled 'Standards Track'. It displays the document's structure, including sections like 'Introduction', 'The SMTP Model', and 'Simple Mail Transfer Protocol'. On the right is a screenshot of a code editor showing the generated Java code for the SMTP protocol. The code is written in a style that follows the protocol specification, using annotations like 'aux global protocol' and 'rec' to define session types. Below the code editor is a terminal window showing the output of a command-line application named 'Smtp.java', which includes status messages like 'Checking states: 50' and 'Checked all states: 67'.

- ▶ Homepage and tutorials <http://www.scribble.org/>
- ▶ Scribble-Java source <https://github.com/scribble/scribble-java>

Scribble: protocol-specific API generation from session types

- ▶ MPST-based toolchain for Java (also .NET, Go, ...)
 - ▶ Protocol validation: MPST syntax + model checking (+ SAT solving)
 - ▶ Endpoint programming using generated distributed APIs

The screenshot shows an IDE interface with two tabs: "Smtp.scr" and "SmtpC.java". The "Smtp.scr" tab contains Scribble protocol code, specifically a sequence of socket operations (send, receive) for sending an email. The "SmtpC.java" tab shows the generated Java code, which is a class with methods corresponding to the protocol steps. Below the tabs is a code editor window displaying the generated Java code. At the bottom of the IDE is a toolbar with icons for Problems, Javadoc, Declaration, Search, Console, and JUnit. A status bar indicates "0 errors, 59 warnings, 0 others". A dropdown menu is open, showing "Description" and "Warnings (59 items)".

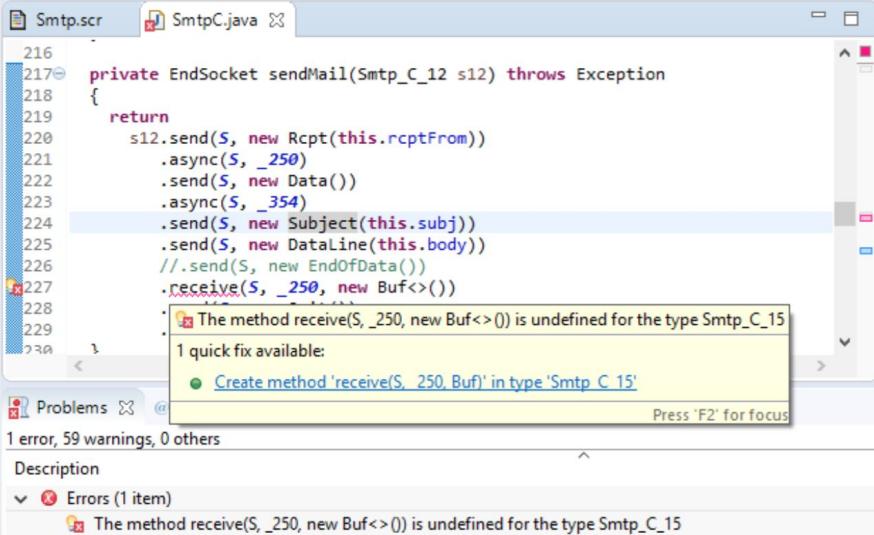
```
216
217     private EndSocket sendMail(Smtp_C_12 s12) throws Exception
218     {
219         return
220             s12.send($, new Rcpt(this.rcptFrom))
221                 .async($, _250)
222                 .send($, new Data())
223                 .async($, _354)
224                 .send($, new Subject(this.subj))
225                 .send($, new Dataline(this.body))
226                 .send($, new EndOfData())
227                 .receive($, _250, new Buf<>())
228                 .send($, new Quit())
229                 .receive($, _221, new Buf<>());
230     }

```

- ▶ Homepage and tutorials <http://www.scribble.org/>
- ▶ Scribble-Java source <https://github.com/scribble/scribble-java>

Scribble: protocol-specific API generation from session types

- ▶ MPST-based toolchain for Java (also .NET, Go, ...)
 - ▶ Protocol validation: MPST syntax + model checking (+ SAT solving)
 - ▶ Endpoint programming using generated distributed APIs



The screenshot shows an IDE interface with two tabs: "Smtp.scr" and "SmtpC.java". The "SmtpC.java" tab contains Java code for an SMTP session type. A tooltip is displayed over the line of code: ".receive(S, _250, new Buf<>())". The tooltip message is: "The method receive(S, _250, new Buf<>()) is undefined for the type Smtp_C_15". Below the tooltip, it says "1 quick fix available:" and "Create method 'receive(S, _250, Buf)' in type 'Smtp_C_15'". The "Problems" view at the bottom shows one error: "The method receive(S, _250, new Buf<>()) is undefined for the type Smtp_C_15".

```
216
217  private EndSocket sendMail(Smtp_C_12 s12) throws Exception
218  {
219      return
220      s12.send(S, new Rcpt(this.rcptFrom))
221          .async(S, _250)
222          .send(S, new Data())
223          .async(S, _354)
224          .send(S, new Subject(this.subj))
225          .send(S, new DataLine(this.body))
226          // .send(S, new EndOfData())
227          .receive(S, _250, new Buf<>())
228
229
230 }
```

Problems

1 error, 59 warnings, 0 others

Description

Errors (1 item)

The method receive(S, _250, new Buf<>()) is undefined for the type Smtp_C_15

- ▶ Homepage and tutorials <http://www.scribble.org/>
- ▶ Scribble-Java source <https://github.com/scribble/scribble-java>

PL4DS challenges (from a session types view)

- ▶ Context: “high-level” (TCP-like) distributed applications
 - ▶ (*Choreographic*) *specifications* of real-world protocols
 - ▶ Engineering, verification, expressiveness, ...
 - ▶ *Implementations* of endpoint programs in “mainstream” languages
 - ▶ Native, distributed (heterogeneous), interoperability, ...
 - ▶ Practical safety guarantees
 - ▶ I/O primitives vs. libraries, channel linearity/aliasing/..., multithreaded/event-driven/actors/...
- ▶ Cross-discipline approaches to safe distributed programming
 - ▶ Centering on language-based protocol enforcement
 - ▶ Type checking/inference, model checking, SAT solving, code generation, meta programming (e.g., type providers), ...

Mira Mezini

TU Darmstadt



Programming languages:

- interactive data-intensive distributed applications
- software defined networking
- privacy-preserving computations

Program analysis:

- vulnerability/malware detection
- automated software engineering (“big code”)

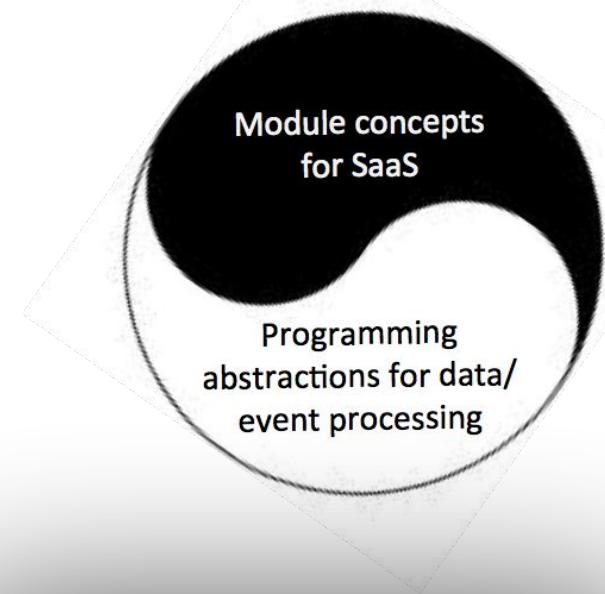
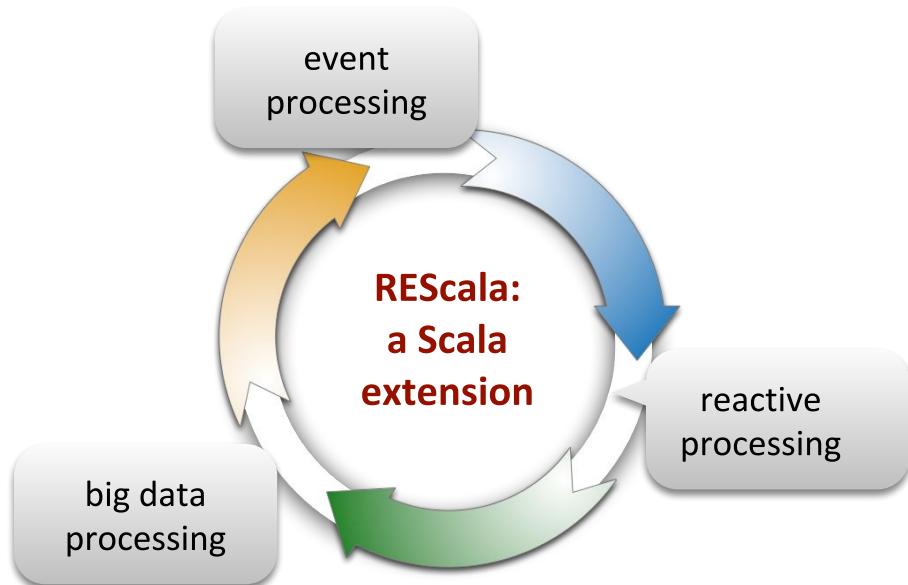


European Research Council

Established by the European Commission

Supporting top researchers
from anywhere in the world

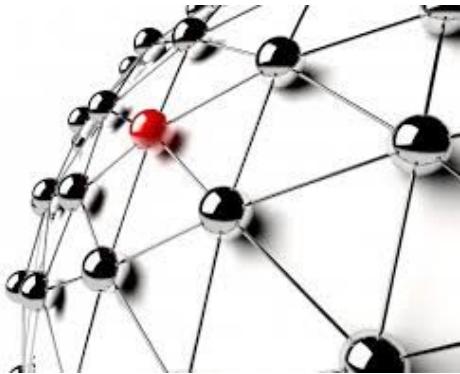
PACE



WORKING HYPOTHESIS

Application design quality would benefit from **integration and unification of abstractions** for reactive, event and big data processing, making them **composable**.

ReScala: Reactive programming for distributed applications (**REScala**)



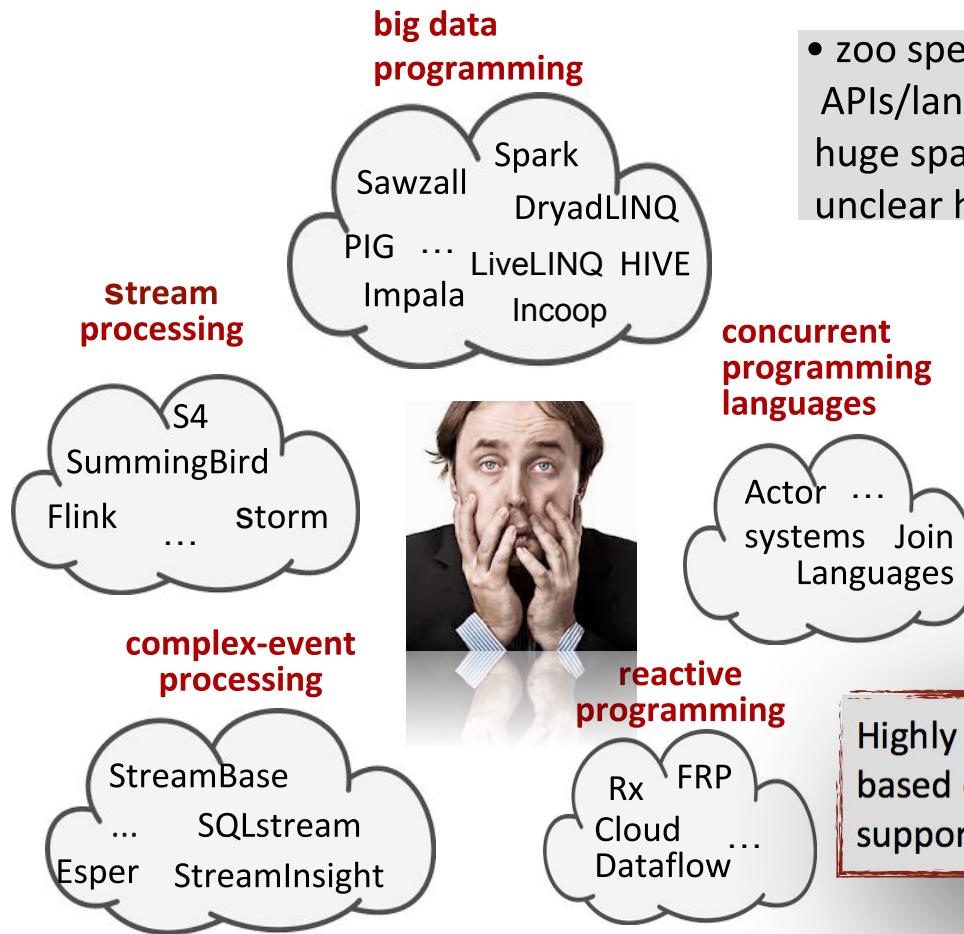
Ubiquitous callbacks!

Hypothesis: Distributed reactives reduce the complexity due to ubiquitous pub-sub systems.

- Distributed reactive values (signals/events)
- Decentralized, concurrent, fault-tolerant reactive updates
- Different levels of consistency (strong consistency via glitch-freedom; eventual consistency via CRDTs ... consistency types)

BUT: Fixed join semantics (“most-recent”)

DATA/EVENT AGGREGATION ...



- zoo specialised data/event processing APIs/languages...
- huge space of feature/semantic variability
- unclear how features across domains compose/interact

Highly needed: Unification/generalization/composability based on common language foundations with systematic support for customisability.

CorrL: Unified language framework for data/even correlations ([CorrL](#))

- Conceptual core: Joins = cartesian product of asynchronous streams
- Custom join semantics expressed via algebraic effects /handlers
- Streamlined, extensible language - correlation features as libraries.
- Correlation semantics can co-exist and even be dynamically selected.

an analogy by Ion Stoica



First cellular phones



specialized devices



Unified device
(smartphone)

Philipp Haller

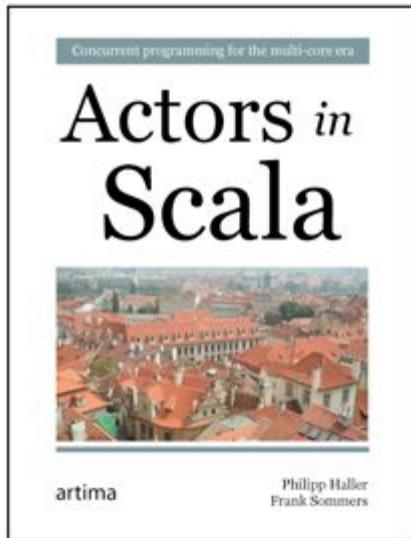
- 2014-present: KTH Royal Institute of Technology
- 2012-2014: Typesafe Inc.
- 2011-2012: Stanford and EPFL
- PhD EPFL, 2010

Research interests:

- Reactive and distributed programming
- Type systems for concurrency and distribution
- Deterministic concurrency
- Weak consistency



Reactive programming and actors



- Programming languages for reactive systems
- Actors
- Foundations
- Experimental evaluation
- Geo distribution

Philipp Haller, Frank Sommers
Artima Press, 2011
ISBN 978-0-9815316-5-6

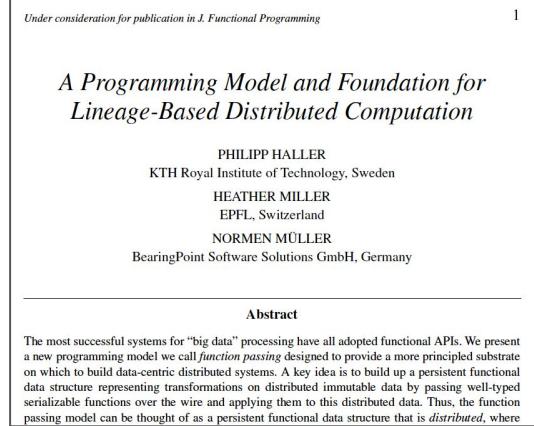
Linear and affine types

- Linear and affine types "for the masses":
 - scale to fully-fledged languages
 - applicable to real code
- Typestate
- "Types for concurrency and distribution"



Haller and Loiko.
LaCasa: lightweight affinity and object capabilities in Scala. OOPSLA 2016

Deterministic concurrency



- Programming models
- Foundations
- Distribution
- Fault tolerance

Haller, Miller, and Müller.
A Programming Model and Foundation for
Lineage-Based Distributed Computation.
Journal of Functional Programming, 2018. To appear.

Annette Bieniusa

- Since 2012 Senior researcher / lecturer at TU Kaiserslautern in AG Softech
- 2011-12 Postdoc at Inria Paris / LIP6
- 2007-11 PhD at University of Freiburg (with Peter Thiemann)



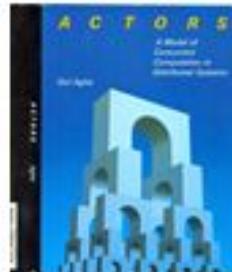
Research interests

- Shared state on highly-scalable distributed systems
 - CRDTs
 - (Distributed) Software Transactional memory
- Developing correct applications on weakly-consistent data stores
- Framework for offline functionality in mobile apps



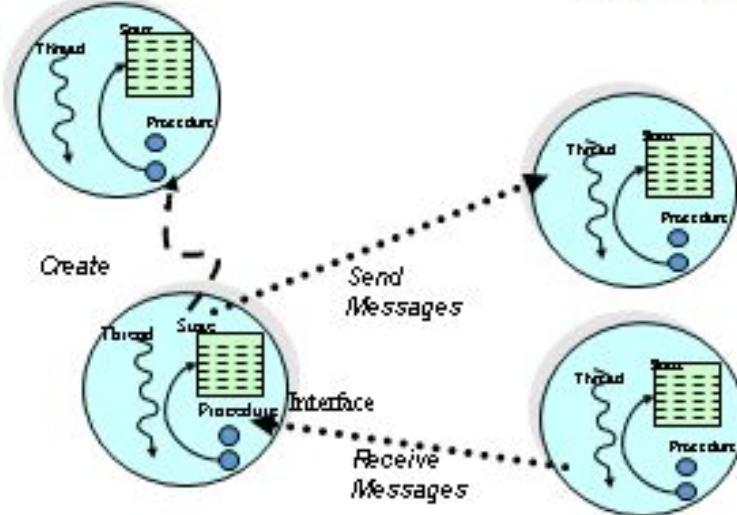
- Geo-replicated data store with low latency and high availability
- Data model: CRDTs
- Transactional Causal+ Consistency
- Transactions group updates and reads from snapshot
- Specification, verification and testing with Repliss
 - Data invariants + Invariants over history of invocations
- Event-based Parallel Temporal Logic (EPTL)

```
val chat: CrdtMap_dw[ChatId, CrdtSet_rw[MessageId]]  
val message: CrdtMap_dw[MessageId, {  
    author: CrdtMvReg[UserId],  
    content: CrdtMvReg[String],  
    chat: CrdtMvReg[ChatId],  
}]  
  
def sendMessage(from: UserId,  
               content: String, to: ChatId) {  
    atomic {  
        val m = new MessageId  
        message(m).author.assign(from)  
        message(m).content.assign(content)  
        message(m).chat.assign(to)  
        chat(to).add(m)  
    }  
}  
  
def editMessage(id: MessageId, newContent: String) {  
    atomic {  
        if (messages(id).exists) {  
            message(id).content.assign(newContent)  
        }  
    }  
}  
  
def deleteMessage(id: MessageId) {  
    atomic {  
        val c = first(message(id).chat)  
        chat(c).remove(id)  
        message.delete(id)  
    }  
}
```



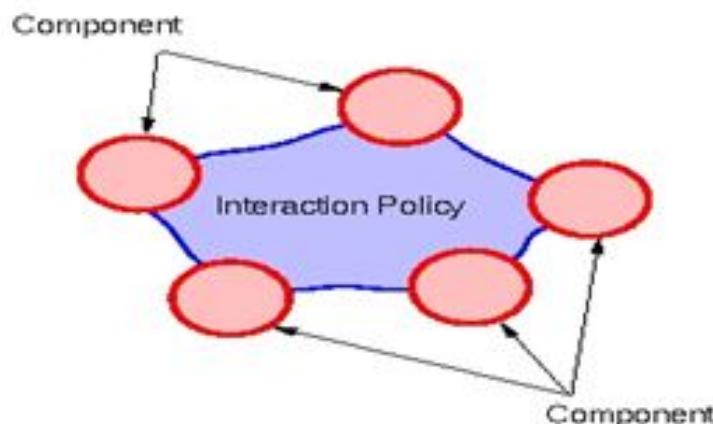
The Actor Model

- The world is naturally decomposed into autonomous objects that act concurrently
- There is no action at a distance: asynchronous communication
- Scalable, reactive and messaging applications (Twitter, LinkedIn, Facebook Chat, ...)
- Database and Cloud applications (Orleans)



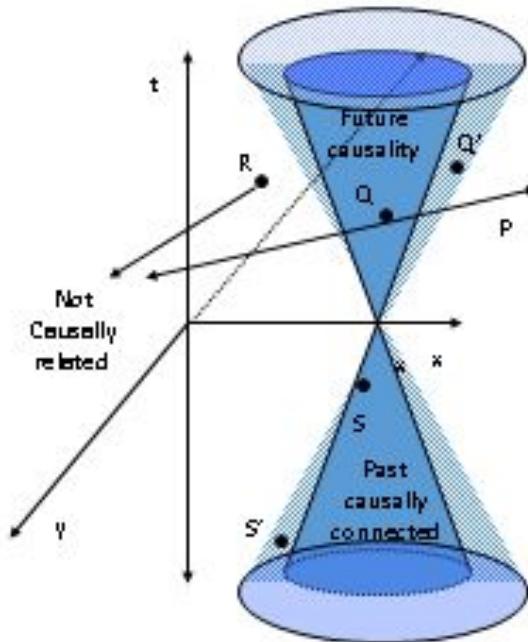
Abstracting Interaction

- How to express *coordination* (e.g. *synchronization intention*) modularly?
- Separate *interaction policy* from *protocols* used to implement policy
 - Multiparty Session types
 - Synchronizers
 - Streams, joins, ..
 - Workflows



Actors Sample points in Space-Time

- Events separated in space are separated in time:
 - Scheduling delays
 - Latency and communication delays
- Such delays are probabilistic in nature
- Probabilistic cone
 - Probabilistic Actor Programming
 - Statistical Model Checking



Cyberphysical Space

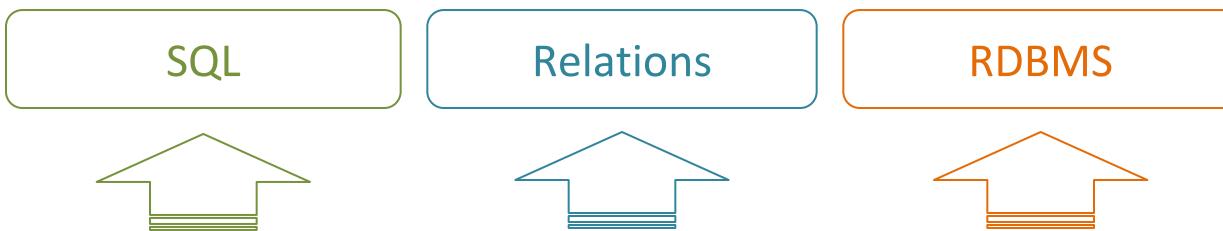


- **Continuous as well as discrete variables as first-class objects**
 - Location, time, motion, strain, vibration, etc. are types of physical attributes best represented by continuous variables
- **Stochastic variables as first-class objects**
 - Real-world systems (IoT, sensor networks, web search, robots) involve probabilistic analysis and statistical decisions
- **Interoperation of suitable programming abstractions:**
 - Actors
 - Constraints

Declarative Data Processing and Mosaics

A DBMS and Big Data Processing Perspective

A Billion \$\$\$ Mantra...



Declarative Data Processing

An effective, formal foundation based on relational algebra and calculus (Codd '71).

A simple, high-level language for querying data (Chamberlin '74).

An efficient, low-level execution environment tailored towards the data (Selinger '79).

With 40+ years of success...



SQL

Relations

RDBMS



Declarative Data Processing

Is Being Revised



SQL

Relations

RDBMS



Declarative Data Processing



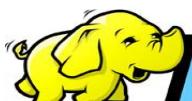
Second-Order
Functions

Distributed
Collections

Parallel Dataflow
Engines



Flink

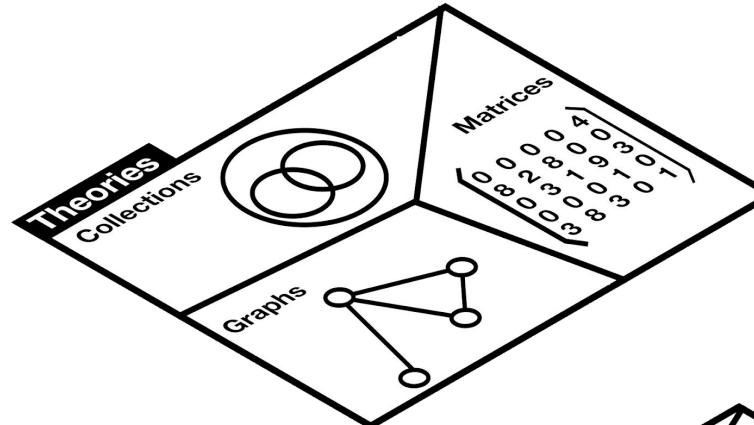


hadoop

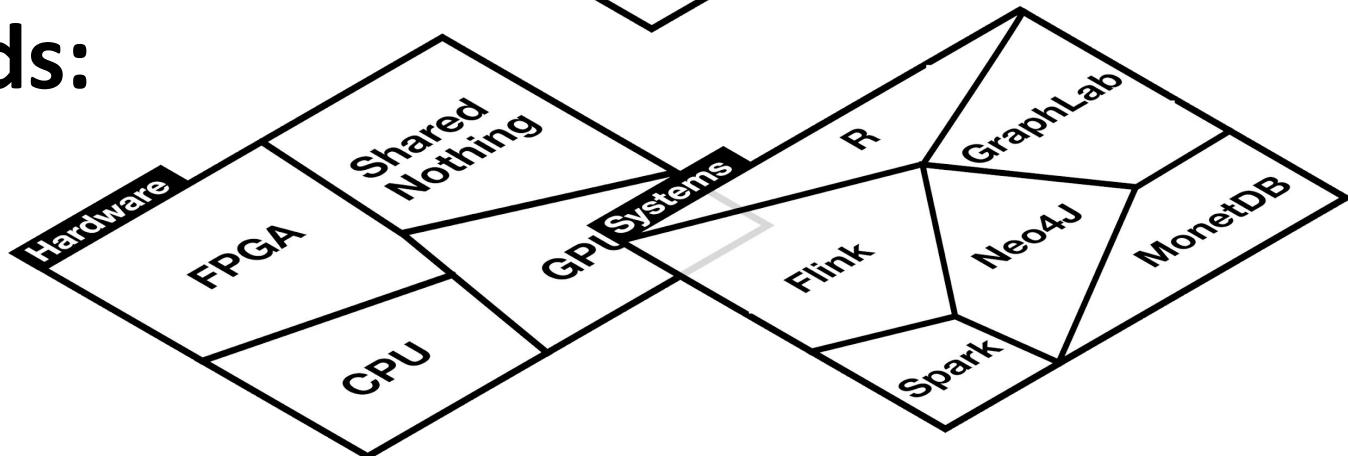
Spark

Mosaics

Frontend:



Backends:



Research

1. Unifying Modelling Across Theories
2. Cross Theory Optimization
3. Optimizing Across Engines
4. Predicting and Learning Program Runtimes
5. Optimizing Across Hardware
6. Generating Hardware-Targeted Code

